



SIMON FRASER UNIVERSITY
ENGAGING THE WORLD

CMPT 431 Distributed Systems

Fall 2019

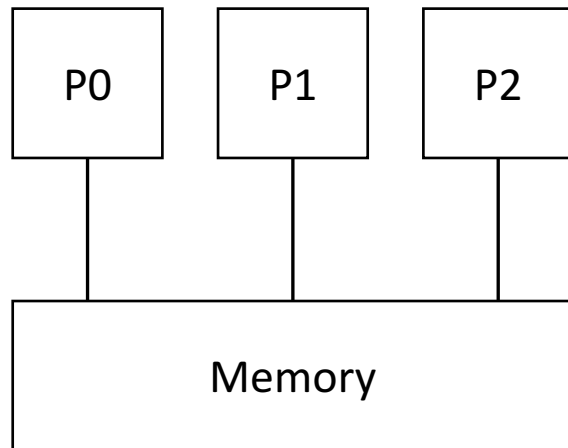
Parallel Computing

<https://www.cs.sfu.ca/~keval/teaching/cmpt431/fall19/>

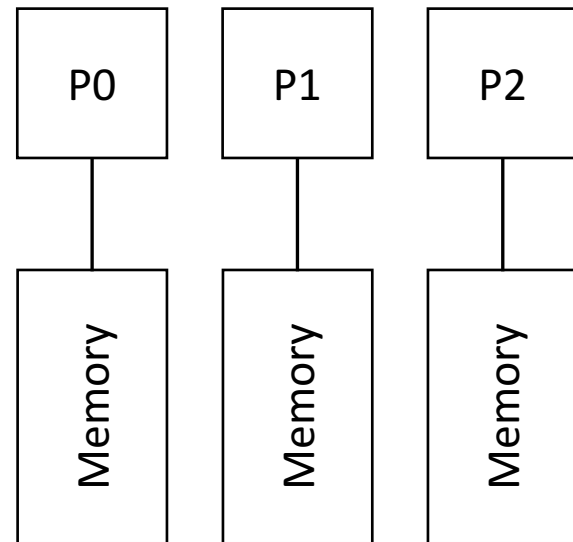
Instructor: Keval Vora

Memory Architectures

Shared Memory



Distributed Memory



Shared Memory Architectures

- Processors access memory as a global address space
- Memory updates by one processor are visible to others
- Easier to program with global address space
- Typically fast memory access (supported by hardware)
- Difficult to scale
- Adding CPUs (geometrically) increases traffic
- Need to synchronization memory accesses

Shared Memory: Race Condition

```
withdraw(int account_id, int amount) {  
    balance = get_balance(account_id);  
    if(balance >= amount) {  
        set_balance(account_id,  
                     balance - amount);  
        eject(amount);  
    }  
}
```

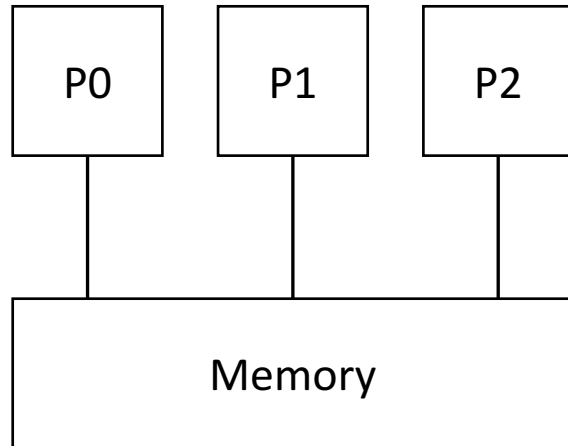
Distributed Memory Architectures

- Nodes (processors + memories) connected via communication network
- Access to another processor's data via communication protocols (e.g., send-receive calls)
- Scalable (both processor and memory)
- Local access is fast (no coordination required)
- Cost effective
- Difficult to program
- Communication/Synchronization is difficult to manage

Shared Memory: UMA & NUMA

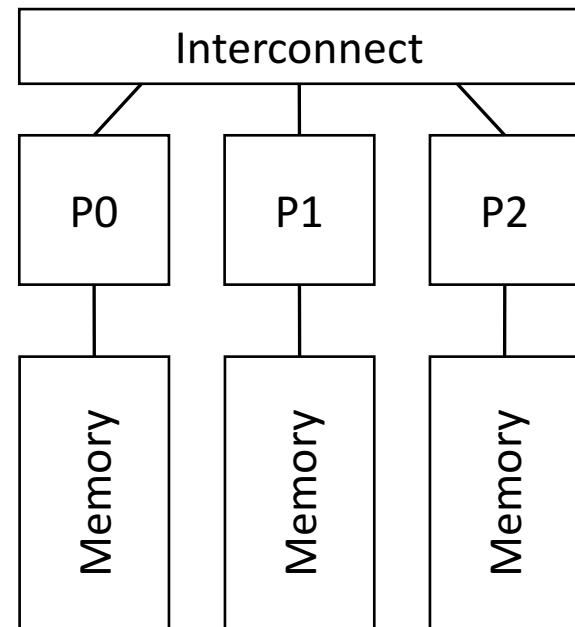
Uniform Memory Access (UMA)

- Typically Symmetric Multiprocessors (SMP)
- **Equal** memory access time



Non-Uniform Memory Access (NUMA)

- Typically multiple SMPs that can access each other's memories
- Memory access times are **not equal**



Communication Models

- Shared Memory

- Tasks share a **common address space** they access asynchronously
- Mutexes/Locks used to **control access** to shared memory
- Compiler translates variables into global memory addresses

- Message Passing

- **No shared** address space
- Tasks use their own **local memories**
- Data transfer often requires coordination: receive matching send

Parallel Programming Models

- Data parallel
 - Parallel operations over a collection of data items
 - Tasks collectively work on the collection and each task works on a different partition (subset) of the collection
 - E.g., increment all elements in an array by 1 (**map operation**)
- Task Parallel
 - Tasks defined based on the operations to be performed
 - Each task performs an operation which is **different** from that performed by other task
 - Operations should be **safe** to be concurrently executable with other operations
 - May or may not operate on the same collection of data items
 - E.g., pipelining

Parallel Programming

- Decompose problem into sub-problems
- Map sub-problems to concurrent tasks
- Ensure **dependencies** are correctly satisfied
 - E.g., Op1 in thread 1 should happen before op2 in thread 2
 - Coordination via synchronization/communication
- Overall execution is a mix of parallel and sequential executions

Measuring Performance

- How fast does a job compute?
 - Elapsed time (**Latency**)
 - compute + communicate + synchronize
- How many jobs complete in a given time?
 - **Throughput**
- How well does the system scale?
 - Increasing *processors* (compute nodes)
 - Increasing problem size (constant work per processor)

Weak v/s strong scaling: https://en.wikipedia.org/wiki/Scalability#Weak_versus_strong_scaling

Performance Metrics

- Speedup, $S_p = \frac{\text{Execution time using 1 processor system } (T_1)}{\text{Execution time using p processor system } (T_p)}$
- Efficiency = $\frac{S_p}{p}$
- Cost, $C_p = p \times T_p$ Optimal if $C_p = T_1$

Amdahl's Law

- f = fraction of problem that is sequential
 - $(1 - f)$ = fraction of problem that is parallel
- Fastest parallel time, $T_p = T_1 \times (f + \frac{1-f}{p})$
- Speedup with p processors, $S_p = \frac{1}{f + \frac{(1-f)}{p}}$

Amdahl's Law

- Gives an upper bound on speedup
- Only fraction $(1 - f)$ is shared by p processors
 - Increasing p cannot speed up fraction f
- Speedup with p processors, $S_p = \frac{1}{f + \frac{(1-f)}{p}}$
- Upper bound on speedup at $S_\infty = \frac{1}{f}$
- Example: $f = 2\%$, $S_\infty = 1 / 0.02 = 50$

Amdahl's Law

$$S_p = \frac{1}{f + \frac{(1-f)}{p}}$$

