



SIMON FRASER UNIVERSITY
ENGAGING THE WORLD

CMPT 431 Distributed Systems

Fall 2019

Parallel Programming Examples

<https://www.cs.sfu.ca/~keval/teaching/cmpt431/fall19/>

Instructor: Keval Vora


Parallel Programming

- Goal: Parallelize some simple sequential programs
- Understand some of the challenges involved
- Not tied to any specific language
- Homework
 - Implement each parallel program and check performance
 - Try to improve performance with better parallelization strategies

Monte Carlo π Estimation

```
circle_count = 0;
for (uint i = 0; i < n; i++) {
    x = get_random(0, 1);
    y = get_random(0, 1);
    if ((x, y) inside circle)
        ++circle_count;
}
pi_value = 4.0 * circle_count / n;
```


$$\frac{A(circle)}{A(square)} = \frac{\pi r^2}{4r^2}$$



Monte Carlo π Estimation

```
circle_count = 0;  
create T threads  
for each thread in parallel {  
    for (uint i = 0;  
        i < approx(n/T); i++) {  
        x = get_random(0, 1);  
        y = get_random(0, 1);  
        if ((x, y) inside circle)  
            ++circle_count;  
    }  
}  
pi_value = 4.0 * circle_count / n;
```


$$\frac{A(circle)}{A(square)} = \frac{\pi r^2}{4r^2}$$



Monte Carlo π Estimation

$$\frac{A(circle)}{A(square)} = \frac{\pi r^2}{4r^2}$$


```
circle_count = 0;  
create T threads  
for each thread in parallel {  
    for (uint i = 0;  
        i < approx(n/T); i++) {  
        x = get_random(0, 1);  
        y = get_random(0, 1);  
        if ((x, y) inside circle)  
            ++circle_count;  
    }  
}  
pi_value = 4.0 * circle_count / n;
```



Monte Carlo π Estimation

$$\frac{A(circle)}{A(square)} = \frac{\pi r^2}{4r^2}$$


```
circle_count = 0;
create T threads
for each thread in parallel {
    for (uint i = 0;
        i < approx(n/T); i++) {
        x = get_random(0, 1);
        y = get_random(0, 1);
        if ((x, y) inside circle)
            lock();
            ++circle_count;
            unlock();
        }
    }
    pi_value = 4.0 * circle_count / n;
```



Monte Carlo π Estimation

$$\frac{A(circle)}{A(square)} = \frac{\pi r^2}{4r^2}$$


```
circle_count = 0;
create T threads
for each thread in parallel {
    for (uint i = 0;
        i < approx(n/T); i++) {
        x = get_random(0, 1);
        y = get_random(0, 1);
        if ((x, y) inside circle)
            lock();
            ++circle_count;
            unlock();
    }
}
pi_value = 4.0 * circle_count / n;
```



serial code

Monte Carlo π Estimation


$$\frac{A(circle)}{A(square)} = \frac{\pi r^2}{4r^2}$$



```
circle_count = 0;
create T threads
for each thread in parallel {
    local_circle_count = 0;
    for (uint i = 0;
        i < approx(n/T); i++) {
        x = get_random(0, 1);
        y = get_random(0, 1);
        if ((x, y) inside circle)
            ++local_circle_count;
    }
    lock();
    circle_count += local_circle_count;
    unlock();
}
pi_value = 4.0 * circle_count / n;
```


Monte Carlo π Estimation

$$\frac{A(circle)}{A(square)} = \frac{\pi r^2}{4r^2}$$




```
circle_count = 0;
create T threads
for each thread in parallel {
    local_circle_count = 0;
    for (uint i = 0;
        i < approx(n/T); i++) {
        x = get_random(0, 1);
        y = get_random(0, 1);
        if ((x, y) inside circle)
            ++local_circle_count;
    }
    lock();
    circle_count += local_circle_count;
    unlock();
}
pi_value = 4.0 * circle_count / n;
```

use atomics
to eliminate
locks

Monte Carlo π Estimation

$$\frac{A(circle)}{A(square)} = \frac{\pi r^2}{4r^2}$$



```
circle_count = 0;
create T threads
for each thread in parallel {
    local_circle_count = 0;
    for (uint i = 0;
        i < approx(n/T); i++) {
        x = get_random(0, 1);
        y = get_random(0, 1);
        if ((x, y) inside circle)
            ++local_circle_count;
    }
    atomic_add(circle_count, local_circle_count);
}
pi_value = 4.0 * circle_count / n;
```

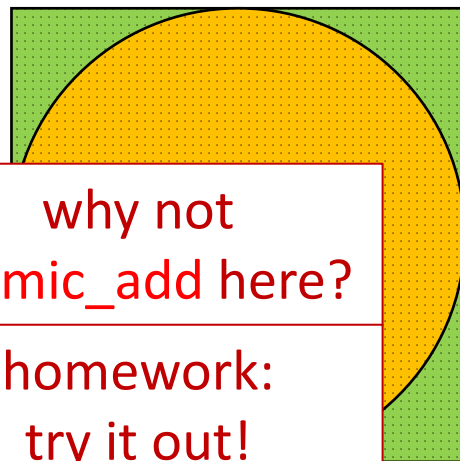
check out
FAA, CAS

CAS: <https://en.wikipedia.org/wiki/Compare-and-swap>
FAA: <https://en.wikipedia.org/wiki/Fetch-and-add>
C++11 std::atomic: <https://en.cppreference.com/w/cpp/atomic>

Monte Carlo π Estimation

$$\frac{A(circle)}{A(square)} = \frac{\pi r^2}{4r^2}$$

```
circle_count = 0;
create T threads
for each thread in parallel {
    local_circle_count = 0;
    for (uint i = 0;
        i < approx(n/T); i++) {
        x = get_random(0, 1);
        y = get_random(0, 1);
        if ((x, y) inside circle)
            ++local_circle_count;
    }
    atomic_add(circle_count, local_circle_count);
}
pi_value = 4.0 * circle_count / n;
```



why not
atomic_add here?

homework:
try it out!


CAS: <https://en.wikipedia.org/wiki/Compare-and-swap>
FAA: <https://en.wikipedia.org/wiki/Fetch-and-add>
C++11 std::atomic: <https://en.cppreference.com/w/cpp/atomic>

Triangle Counting


```
triangles(u, v) {  
    return |outNeighbors(v)  $\cap$  inNeighbors(u)|;  
}
```

```
count = 0;  
for u in V {  
    for v in outNeighbors(u) {  
        count += triangles(u, v);  
    }  
}  
count = count / 3;
```


parallelize



Triangle



$G = (V, E)$




Triangles

Triangle Counting


```
triangles(u, v) {  
    return |outNeighbors(v)  $\cap$  inNeighbors(u)|;  
}
```

```
count = 0;  
for u in V {  
    for v in outNeighbors(u) {  
        count += triangles(u, v);  
    }  
}  
count = count / 3;
```


write shared



Triangle




$G = (V, E)$




Triangles

Triangle Counting

```
triangles(u, v) {  
    return |outNeighbors(v)  $\cap$  inNeighbors(u)|;  
}  
  
count = 0;  
create T threads  
for each thread in parallel {  
    local_count = 0;  
    for u in subset(V, tid) {  
        for v in outNeighbors(u) {  
            local_count += triangles(u, v);  
        }  
    }  
    atomic_add(count, local_count);  
}  
count = count / 3;
```




Triangle



solution
similar to
 π estimation

PageRank

$$PR[V] = 0.15 + 0.85 \times \sum_{u \in inNeighbors(v)} \frac{PR[u]}{degree(u)}$$



- Used by Google to rank webpages
- Estimates importance of a webpage by counting the number and quality of links to the page
- Iterative algorithm
 - Compute ranks over and over again to account for changes in incoming links' qualities
 - PageRank values stop changing over time

PageRank: <https://en.wikipedia.org/wiki/PageRank>

PageRank

		Vertices						
		0	1	2	3	4	5	...
0								
1								

```
#define CURR i % 2
#define NEXT (i+1) % 2
```

```
for (i=0; i<k; ++i) {
    for v in V {
        sum = 0;
        for u in inNeighbors(v) {
            sum += pagerank[CURR][u] / out_degree(u);
        }
        pagerank[NEXT][v] = 0.15 + 0.85 x sum;
    }
}
```

parallelize

PageRank

		Vertices						
		0	1	2	3	4	5	...
0								
1								

```
#define CURR i % 2
#define NEXT (i+1) % 2
```

```
for (i=0; i<k; ++i) {
    for v in V {
        sum = 0;
        for u in inNeighbors(v) {
            sum += pagerank[CURR][u] / out_degree(u);
        }
        pagerank[NEXT][v] = 0.15 + 0.85 x sum;
    }
}
```

shared data?

separate read and
write locations

PageRank

		Vertices						
		0	1	2	3	4	5	...
0								
1								

```
#define CURR i % 2
#define NEXT (i+1) % 2
```

create T threads

```
for each thread in parallel {
    for (i=0; i<k; ++i) {
        for v in subset(V, tid) {
            sum = 0;
            for u in inNeighbors(v) {
                sum += pagerank[CURR][u] / out_degree(u);
            }
            pagerank[NEXT][v] = 0.15 + 0.85 x sum;
        }
    }
}
```


read and write locations
still different?

PageRank

```
#define CURR i % 2
#define NEXT (i+1) % 2
```

create T threads

```
for each thread in parallel {
    for (i=0; i<k; ++i) {
        for v in subset(V, tid) {
            sum = 0;
            for u in inNeighbors(v) {
                sum += pagerank[CURR][u] / out_degree(u);
            }
            pagerank[NEXT][v] = 0.15 + 0.85 x sum;
        }
    }
}
```




read and write locations
still different?

PageRank

```
#define CURR i % 2
#define NEXT (i+1) % 2
```

create T threads

```
for each thread in parallel {
    for (i=0; i<k; ++i) {
        for v in subset(V, tid) {
            sum = 0;
            for u in inNeighbors(v) {
                lock();
                sum += pagerank[CURR][u] / out_degree(u);
                unlock();
            }
            lock();
            pagerank[NEXT][v] = 0.15 + 0.85 x sum;
            unlock();
        }
    }
}
```



locks don't control
thread progress


PageRank

```
#define CURR i % 2
#define NEXT (i+1) % 2
```

create T threads

```
for each thread in parallel {
    for (i=0; i<k; ++i) {
        for v in subset(V, tid) {
            sum = 0;
            for u in inNeighbors(v) {
                sum += pagerank[CURR][u] / out_degree(u);
            }
            pagerank[NEXT][v] = 0.15 + 0.85 x sum;
        }
        barrier();
    }
}
```


blocks threads until all
threads arrive here



Barrier: [https://en.wikipedia.org/wiki/Barrier_\(computer_science\)](https://en.wikipedia.org/wiki/Barrier_(computer_science))


Dijkstra's SSSP Algorithm

```
pq.enqueue(<0, source>);    // Priority Queue
while(!pq.empty()) {
    u = pq.dequeue().second;
    for v in outNeighbors(u) {
        if(d[u] + w(u, v) < d[v]) {
            d[v] = d[u] + w(u, v);
            pq.enqueue(<d[v], v>);
        }
    }
}
```



Dijkstra's SSSP Algorithm

```
pq.enqueue(<0, source>);    // Priority Queue
while(!pq.empty()) {
    u = pq.dequeue().second;
    for v in outNeighbors(u) {
        if(d[u] + w(u, v) < d[v]) {
            d[v] = d[u] + w(u, v);
            pq.enqueue(<d[v], v>);
        }
    }
}
```



A red rectangular box labeled "shared data" has three red arrows pointing to the variables `d[u]`, `d[v]`, and `w(u, v)` in the code snippet, indicating that these variables represent shared data in a distributed context.

Dijkstra's SSSP Algorithm

Note: concurrent data structures without locks later in this course

```
pq.enqueue(<0, source>);
while(!pq.empty()) {
    u = pq.dequeue().second;
    for v in outNeighbors(u) {
        if(d[u] + w(u, v) < d[v]) {
            d[v] = d[u] + w(u, v);
            pq.enqueue(<d[v], v>);
        }
    }
}
```

```
enqueue() {
    lock();
    . . .
    unlock();
}

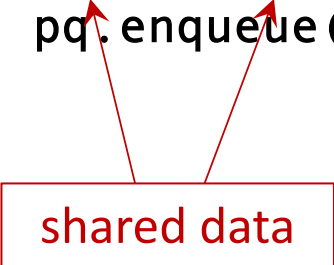
dequeue() {
    lock();
    . . .
    unlock();
}

empty() {
    lock();
    . . .
    unlock();
}
```


Dijkstra's SSSP Algorithm

Note: concurrent data structures without locks later in this course

```
pq.enqueue(<0, source>);
while(!pq.empty()) {
    u = pq.dequeue().second;
    for v in outNeighbors(u) {
        if(d[u] + w(u, v) < d[v]) {
            d[v] = d[u] + w(u, v);
            pq.enqueue(<d[v], v>);
        }
    }
}
```



shared data

- Violating “order” is okay for SSSP because intermediate values are not necessarily incorrect
- However, no “update” should be missed
- $d[\cdot]$ values will get refined and corrected out
- Hence, don't need to order threads (i.e., no barrier)
- This is different from PageRank

```
enqueue() {
    lock();
    . . .
    unlock();
}

dequeue() {
    lock();
    . . .
    unlock();
}

empty() {
    lock();
    . . .
    unlock();
}
```

Dijkstra's SSSP Algorithm

```
pq.enqueue(<0, source>);
```

```
create T threads
```

```
for each thread in parallel {
```

```
    while(!pq.empty()) {
```

```
        u = pq.dequeue().second;
```

```
        if(u == NULL) continue;
```

```
        for v in outNeighbors(u) {
```

```
            if(atomic_min(d[v],
```

```
                        d[u] + w(u, v))) {
```

```
                pq.enqueue(<d[v], v>);
```

```
            }
```

```
        }
```

```
    }
```

```
}
```

to safeguard
against race

- returns true if $d[v]$ got updated as $d[u] + w(u, v)$
- returns false if $d[v]$ already has value $\leq d[u] + w(u, v)$
- can be implemented using CAS (homework)

Dijkstra's SSSP Algorithm

```
pq.enqueue(<0, source>);  
create T threads  
for each thread in parallel {  
    while(!pq.empty()) {  
        u = pq.dequeue().second;  
        if(u == NULL) continue;  
        for v in outNeighbors(u) {  
            if(atomic_min(d[v],  
                           d[u] + w(u, v))) {  
                pq.enqueue(<d[v], v>);  
            }  
        }  
    }  
}
```

What about
termination?

thread 1 gets
pq.empty() = true,
but thread 2 calls
pq.enqueue() just after that

adding barrier()
will not help
(why?)

Dijkstra's SSSP Algorithm

```
done = 0;
pq.enqueue(<0, source>);
create T threads
for each thread in parallel {
  while(true) {
    if(pq.empty()) {
      atomic_add(done, 1);
      while(pq.empty()) {
        if(done == T) break;
      }
      if(done == T) break;
      atomic_add(done, -1);
      continue;
    }
  }
}
```

termination logic

```
u = pq.dequeue().second;
if(u == NULL) continue;
for v in outNeighbors(u) {
  if(atomic_min(d[v],
               d[u] + w(u, v))) {
    pq.enqueue(<d[v], v>);
  }
} // end of while
} // end of for-thread
```

Parallel Algorithm Design

- Identifying independent tasks is not enough
 - Expressing the independent tasks so that they collaboratively accomplish the goal can be challenging
 - Often involves restructuring algorithm
- Correctness based on algorithm needs
 - E.g., ordering threads in PageRank v/s SSSP
- Sub-components need to be made thread-safe
 - E.g., priority queue in SSSP
- Termination detection
 - Especially important when logic lacks synchrony (e.g., no barriers)
- Synchronization
 - locks, barriers
 - Atomics for simple operations are usually much faster

Homework

- Implement each parallel program and check performance
- Try to improve performance with better parallelization strategies

Reading (for Next Class)



- [AMP] Chapter 3
 - Upto 3.7
- [Paper] Linearizability: A Correctness Condition for Concurrent Objects: <https://cs.brown.edu/~mph/HerlihyW90/p463-herlihy.pdf>
 - Upto section 3
- [Paper] How to Make a Multiprocessor Computer That Correctly Executes Multiprocess Programs: <https://www.microsoft.com/en-us/research/uploads/prod/2016/12/How-to-Make-a-Multiprocessor-Computer-That-Correctly-Executes-Multiprocess-Programs.pdf>