



SIMON FRASER UNIVERSITY
ENGAGING THE WORLD

CMPT 431 Distributed Systems

Fall 2019

Reasoning Correctness

<https://www.cs.sfu.ca/~keval/teaching/cmpt431/fall19/>

Instructor: Keval Vora

Concurrent Programs


- How to design concurrent programs?
- What does concurrency mean when **ordering** is important?
 - E.g., concurrent queues, concurrent stacks
- How to verify whether concurrent solution is correct?
 - Naïve: Check for all possible results. Infeasible.
- Properties that can help in:
 - developing concurrent solutions
 - analyzing concurrent solutions
 - provide correctness and progress guarantees

Reading



- [AMP] Chapter 3
 - Upto 3.7
- [Paper] Linearizability: A Correctness Condition for Concurrent Objects: <https://cs.brown.edu/~mph/HerlihyW90/p463-herlihy.pdf>
 - Upto section 3
- [Paper] How to Make a Multiprocessor Computer That Correctly Executes Multiprocess Programs: <https://www.microsoft.com/en-us/research/uploads/prod/2016/12/How-to-Make-a-Multiprocessor-Computer-That-Correctly-Executes-Multiprocess-Programs.pdf>


Single-Enqueuer/Single-Dequeuer Queue



- Sequential code is available

```
1  class Queue<T> {
2      volatile int head = 0, tail = 0;
3      T[] items;
4      public Queue(int capacity) {
5          items = (T[])new Object[capacity];
6          head = 0; tail = 0;
7      }
8      public void enq(T x) throws FullException {
9          if (tail - head == items.length)
10             throw new FullException();
11          items[tail % items.length] = x;
12          tail++;
13      }
14      public T deq() throws EmptyException {
15          if (tail - head == 0)
16             throw new EmptyException();
17          T x = items[head % items.length];
18          head++;
19          return x;
20      }
21  }
```

Single-Enqueuer/Single-Dequeuer Queue




- Concurrent code is same!
 - How is it correct?
 - `enq()` & `deq()` can happen at same time?

```
1  class WaitFreeQueue<T> {
2      volatile int head = 0, tail = 0;
3      T[] items;
4      public WaitFreeQueue(int capacity) {
5          items = (T[])new Object[capacity];
6          head = 0; tail = 0;
7      }
8      public void enq(T x) throws FullException {
9          if (tail - head == items.length)
10             throw new FullException();
11          items[tail % items.length] = x;
12          tail++;
13      }
14      public T deq() throws EmptyException {
15          if (tail - head == 0)
16             throw new EmptyException();
17          T x = items[head % items.length];
18          head++;
19          return x;
20      }
21  }
```

Lock-based Queue

```
1  class LockBasedQueue<T> {
2      int head, tail;
3      T[] items;
4      Lock lock;
5      public LockBasedQueue(int capacity) {
6          head = 0; tail = 0;
7          lock = new ReentrantLock();
8          items = (T[])new Object[capacity];
9      }
10     public void enq(T x) throws FullException {
11         lock.lock();
12         try {
13             if (tail - head == items.length)
14                 throw new FullException();
15             items[tail % items.length] = x;
16             tail++;
17         } finally {
18             lock.unlock();
19         }
20     }
```



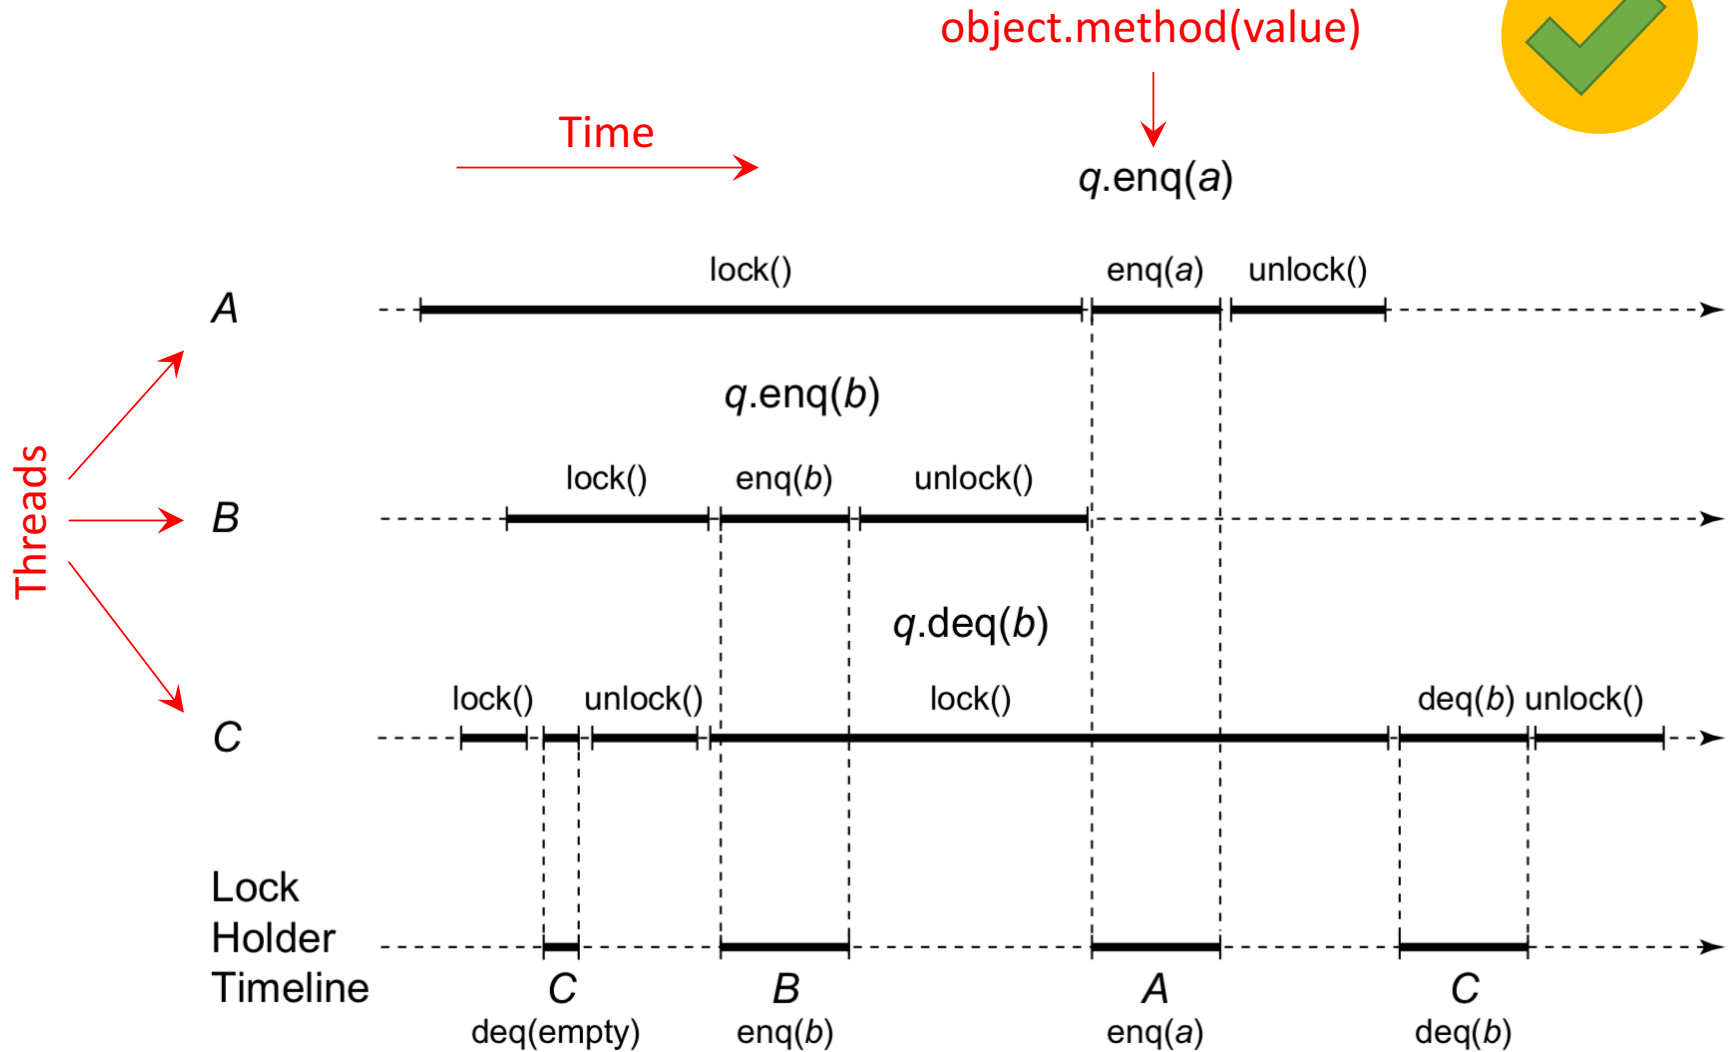
Easy to reason correctness:
no two operations happen at the `same time`

```
21     public T deq() throws EmptyException {
22         lock.lock();
23         try {
24             if (tail == head)
25                 throw new EmptyException();
26             T x = items[head % items.length];
27             head++;
28             return x;
29         } finally {
30             lock.unlock();
31         }
32     }
33 }
```

Operations & Execution Histories

- Operations defined based on the context
 - Queue example: `enq()` and `deq()` are primary operations
- Operations are ones that are performed by concurrent tasks, and that we want to analyze
- Plot **Execution Histories** on timelines to visually analyze when operations (can) occur
 - **Execution Histories**: finite sequence of operations

Execution Histories



Execution Histories



E(x) A

D(y) A

E(y) B

operation(value) Thread

E = Enqueue
D = Dequeue



E(x) A

D(y) A

E(z) A

E(y) B

D(x) B



E(x) A

D(y) A

E(y) B

D(y) C


Correctness Conditions

- Defining the set of **allowable histories**
- “allowable” → Depends on the context
- Degree of concurrency directly affected
 - Allow more histories → higher concurrency
- Sequential Consistency
- Linearizability
- (Serializability)


Sequential Executions

- Sequential methods have clear specifications in form of **preconditions** and **postconditions**
- With concurrency, multiple threads might invoke methods at the same time, making it difficult to reason about preconditions and postconditions in **“isolation”**
 - E.g., what does it mean for two overlapping enqueue operations? Which element will be dequeued first?
- Solution: show equivalence of histories

Method calls should appear to happen in a one-at-a-time sequential order



Sequential Consistency




- **[P1]** Method calls should appear to happen in a one-at-a-time sequential order
- **[P2]** Method calls should appear to take effect in **program order** (i.e., order defined by single thread's code)
 - Purely sequential computations behave the way we would expect
- To verify, reorder method calls sequentially so that:
 - they are consistent with program order
 - meet the object's sequential specification

How to Make a Multiprocessor Computer That Correctly Executes Multiprocess Programs, Leslie Lamport, IEEE TC 1979


Sequential Consistency




- Counter-intuitive because it seems to violate FIFO




- Note: concurrent operations are allowed
 - as long as the two properties are satisfied

Sequential Consistency: Compositional?




- Is the sub-history for p sequentially consistent? 
 - Reorder so that `p.enq(y)` comes before `p.enq(x)`
- Is the sub-history for q sequentially consistent? 
 - Reorder so that `q.enq(x)` comes before `q.enq(y)`
- Is the entire history sequentially consistent? 
 - Need to explicitly check because **SC is not compositional**

Sequential Consistency: Not Compositional



- Is the entire history sequentially consistent?
 - p.deq(y) suggests: p.enq(y) \rightarrow p.enq(x) [red arrow]
 - q.deq(x) suggests: q.enq(x) \rightarrow q.enq(y) [blue arrow]
 - Program order [green arrows]
q.enq(y) \rightarrow p.enq(y) and p.enq(x) \rightarrow q.enq(x)
 - Above orderings form a cycle (i.e., cannot be sequential)


Capturing Real-Time Precedence



- Above history is sequentially consistent
- The **real-time precedence** ordering: $q.enq(x) \rightarrow q.enq(y)$
- Intuition
 - Updates should become visible
 - deq operation should return x and not y
- Sequential consistency fails to capture real-time progress

Linearizability

- Captures **real-time precedence** ordering
- Informally: Operations should **take effect instantaneously** between its invocation and response
- Definitions:
 - $\text{inv}(e)$: invocation of event e
 - $\text{res}(e)$: response of event e
 - $e_0 <_H e_1$: $\text{res}(e_0)$ precedes $\text{inv}(e_1)$ in H
- $<_H$ is a partial order
 - Operations unrelated by $<_H$ are said to be concurrent






Linearizability: A Correctness Condition for Concurrent Objects, M. Herlihy and J. Wing, 1990

Linearizability




- **[P1]** Equivalent to some legal sequential history S
(informally same as [P1] for SC)
- **[P2]** $<_H \subseteq <_S$
(informally, real-time precedence ordering should be preserved)
- Note: equivalence between **complete(H')** and S
 - Read text/paper for formalisms

Linearizability



Linearizability

- Read-Write variable



Linearization Points

- **Linearization points** are points where method takes effect
- Useful to verify and design concurrent solutions
- For lock-based implementations, the critical section is its linearization point
- For implementations without any lock, linearization point is typically a single step where **effects of method call become available to other method calls**

Linearization Points

- For implementations without any lock, linearization point is typically a single step where **effects of method call become available to other method calls**
- `enq()`: if full, when exception gets thrown otherwise, when tail gets updated
- `deq()`: if empty, when exception gets thrown otherwise, when head is updated

```
public void enq(T x) throws FullException {  
    if (tail - head == items.length)  
        throw new FullException();  
    items[tail % items.length] = x;  
    tail++;  
}
```

```
public T deq() throws EmptyException {  
    if (tail - head == 0)  
        throw new EmptyException();  
    T x = items[head % items.length];  
    head++;  
    return x;  
}
```

Linearizability


- Every linearizable execution is sequentially consistent, but not vice versa
- Linearizability is composable
 - Formally, H is linearizable if, and only if, for each object x , $H|x$ is linearizable
 - where $H|x$ means sub-history of H containing operations over x
 - Proof in text/paper

Serializability

- Often considered in database transactions
- History is serializable if it is equivalent to a serial execution
- Similar to Linearizability, except doesn't capture real-time
- Data-centric perspective
 - Preserve semantics of transactions
 - Don't care about program order (or threads)

Serializability: <https://en.wikipedia.org/wiki/Serializability>


Serializability



- Intuitively incorrect
 - lock is held by both threads at the same time
- Serializable
 - B's *m.unlock()* can move before A's *m.lock()*
- Not linearizable – Why?
- Typically:
 - Linearizability useful for components like data structures
 - Serializability useful for complex components like transactions

Iterative Processing

- Values are computed repeatedly in a loop




- Example:

- $\text{pagerank}[\text{NEXT}][v] = 0.15 + 0.85 \times \sum_{\{u \in \text{in}(v)\}} \frac{\text{pagerank}[\text{CURR}][u]}{\text{degree}(u)}$
- $\text{sssp}[v] = \min_{\{u \in \text{in}(v)\}} (\text{sssp}[u])$

- Processing semantics defined by algorithm's needs
- **Dependencies** defined based on value flows

Bulk-Synchronous Parallel Semantics


- Stems from Bulk Synchronous Parallel Model
 - Disjoint computation and communication phases separated by barriers



Valiant et al., A Bridging Model for Parallel Computation, CACM 1990.

Bulk-Synchronous Parallel Semantics

- Stems from Bulk Synchronous Parallel Model
 - Disjoint computation and communication phases separated by barriers
- Need to satisfy read-write dependencies
 - Values in iteration i are computed based on values from iteration $i-1$ only
- PageRank example:
 - $\text{pagerank}[\text{NEXT}][v] = 0.15 + 0.85 \times \sum_{\{u \in \text{in}(v)\}} \frac{\text{pagerank}[\text{CURR}][u]}{\text{degree}(u)}$



$$\forall (u, v) \in E, \quad u^t \mapsto v^{t+1}$$

Asynchronous Execution

- Values don't necessarily progress at same rate
- Recall SSSP:
 - $sssp[v] = \min_{\{u \in in(v)\}} sssp[u]$
- Computation is often **data-driven**
- Correctness is typically coupled with algorithmic semantics
 - E.g., latest update must be eventually read
- **Progress & Termination** semantics
 - **Livelock** and **deadlock** conditions need to be verified
 - Termination conditions need to be defined

Progress Conditions

- Deal with performance instead of correctness
- Progress of a system/thread towards completion
- **Blocking Implementations**
 - Delays by one thread can **prevent others from making progress**
 - E.g., Lock-based queue example
- **Non-Blocking Implementations**
 - Delays by one thread **do not prevent others** from making progress
 - E.g., Single-enqueuer/single-dequeuer queue example

Non-Blocking Guarantee

- Wait-Freedom

- Every operation finishes in finite number of steps
- E.g., Single-enqueuer/single-dequeuer queue example
- Strongest non-blocking guarantee, combining guaranteed system-wide throughput and **starvation-freedom**
- Wait-free algorithms can be inefficient

- Lock-Freedom

- Some operation finishes in finite number of steps
- Allows individual threads to starve, but guarantees system-wide throughput
- Useful especially when starvation is unlikely

Non-blocking algorithm: https://en.wikipedia.org/wiki/Non-blocking_algorithm

Reading (for Next Class)



- Introduction to Parallel Computing (Second Edition) by Ananth Grama, Anshul Gupta, George Karypis, Vipin Kumar
 - Chapter 3: Principles of Parallel Algorithm Design
 - Online Access: https://sfu-primo.hosted.exlibrisgroup.com/permalink/f/usv8m3/01SFUL_ALMA51188913690003611