



SIMON FRASER UNIVERSITY  
ENGAGING THE WORLD

# CMPT 431 Distributed Systems

Fall 2019

## Barriers

<https://www.cs.sfu.ca/~keval/teaching/cmpt431/fall19/>

Instructor: Keval Vora

# Reading



- [AMP] Chapter 17
- [Paper] Algorithms for Scalable Synchronization on Shared Memory Multiprocessors:  
<http://web.mit.edu/6.173/www/currentsemester/readings/R06-scalable-synchronization-1991.pdf>

# Barriers

- Synchronization method to pause all threads at a point

```
scary_code();
```

```
barrier();
```

```
really_scary_code();
```

- Fundamental synchronization primitive: must be efficient
- Minimize communication, writes, bus contention, etc.

# Barrier Implementation

Issues?

```
1  public class SimpleBarrier implements Barrier {
2      AtomicInteger count;
3      int size;
4      public SimpleBarrier(int n){
5          count = new AtomicInteger(n);
6          size = n;
7      }
8      public void await() {
9          int position = count.getAndDecrement();
10         if (position == 1) {
11             count.set(size);
12         } else {
13             while (count.get() != 0);
14         }
15     }
16 }
```

Deadlock!  
Not Reusable

# Sense-Reversing Barrier

```
1  public SenseBarrier(int n) {
2      count = new AtomicInteger(n);
3      size = n;
4      sense = false;
5      threadSense = new ThreadLocal<Boolean>() {
6          protected Boolean initialValue() { return !sense; };
7      };
8  }
9  public void await() {
10     boolean mySense = threadSense.get();
11     int position = count.getAndDecrement();
12     if (position == 1) {
13         count.set(size);
14         sense = mySense;
15     } else {
16         while (sense != mySense) {}
17     }
18     threadSense.set(!mySense);
19 }
```

ThreadLocal<Bool>

Can two thread be in different barrier calls?

# Sense-Reversing Barrier

```
1  public SenseBarrier(int n) {
2      count = new AtomicInteger(n);
3      size = n;
4      sense = false;
5      threadSense = new ThreadLocal<Boolean>() {
6          protected Boolean initialValue() { return !sense; };
7      };
8  }
9  public void await() {
10     boolean mySense = threadSense.get();
11     int position = count.getAndDecrement();
12     if (position == 1) {
13         count.set(size);
14         sense = mySense;
15     } else {
16         while (sense != mySense) {}
17     }
18     threadSense.set(!mySense);
19 }
```


Issues?

# Traffic

- Barrier is a fundamental building block – so we should optimize at lower-most level (i.e., reads and writes)
- $O(P)$  traffic on the bus
  - 2P read/write transactions to update **count**
  - 2 write transactions to write **sense** and reset **count**
  - $P-1$  transactions to read updated **sense**
- Serialization on a single shared variable
  - Latency is  $O(P)$
  - **Can we do better?**

# Combining Tree Barrier

- Reduces memory contention by spreading memory accesses across multiple barriers





```

1 public class TreeBarrier implements Barrier {
2     int radix;
3     Node[] leaf;
4     int leaves;
5     ThreadLocal<Boolean> threadSense;
6     public TreeBarrier(int n, int r) {
7         radix = r;
8         leaves = 0;
9         leaf = new Node[n / r];
10        int depth = 0;
11        threadSense = new ThreadLocal<Boolean>() {
12            protected Boolean initialValue() { return true; };
13        };
14        // compute tree depth
15        while (n > 1) {
16            depth++;
17            n = n / r;
18        }
19        Node root = new Node();
20        build(root, depth - 1);
21    }
22    // recursive tree constructor
23    void build(Node parent, int depth) {
24        if (depth == 0) {
25            leaf[leaves++] = parent;
26        } else {
27            for (int i = 0; i < radix; i++) {
28                Node child = new Node(parent);
29                build(child, depth - 1);
30            }
31        }
32    }

```

Issues?

```

    public void await() {
        int me = ThreadID.get();
        Node myLeaf = leaf[me / radix];
        myLeaf.await();
    }

```

```




1 private class Node {
2     AtomicInteger count;
3     Node parent;
4     volatile boolean sense;
5     public Node() {
6         sense = false;
7         parent = null;
8         count = new AtomicInteger(radix);
9     }
10    public Node(Node myParent) {
11        this();
12        parent = myParent;
13    }
14    public void await() {
15        boolean mySense = threadSense.get();
16        int position = count.getAndDecrement();
17        if (position == 1) { // I'm last
18            if (parent != null) { // Am I root?
19                parent.await();
20            }
21            count.set(radix);
22            sense = mySense;
23        } else {
24            while (sense != mySense) {};
25        }
26        threadSense.set(!mySense);
27    }
28 }
29 }

```

# Combining Tree Barriers

- Removed contention on shared variable
  - Count variables separated for each node
- Unpredictable communication
  - Last thread at every level proceeds towards parent
  - Think NUMA architectures

# Static Tree Barriers



```

1 public class StaticTreeBarrier implements Barrier {
2     int radix;
3     boolean sense;
4     Node[] node;
5     ThreadLocal<Boolean> threadSense;
6     int nodes;
7     public StaticTreeBarrier(int size, int myRadix) {
8         radix = myRadix;
9         nodes = 0;
10        node = new Node[size];
11        int depth = 0;
12        while (size > 1) {
13            depth++;
14            size = size / radix;
15        }
16        build(null, depth);
17        sense = false;
18        threadSense = new ThreadLocal<Boolean>() {
19            protected Boolean initialValue() { return !sense; };
20        };
21    }
22    // recursive tree constructor
23    void build(Node parent, int depth) {
24        if (depth == 0) {
25            node[nodes++] = new Node(parent, 0);
26        } else {
27            Node myNode = new Node(parent, radix);
28            node[nodes++] = myNode;
29            for (int i = 0; i < radix; i++) {
30                build(myNode, depth - 1);
31            }
32        }
33    }
34    public void await() {
35        node[ThreadID.get()].await();
36    }
37 }

```

```

1 public Node(Node myParent, int count) {
2     children = count;
3     childCount = new AtomicInteger(count);
4     parent = myParent;
5 }
6 public void await() {
7     boolean mySense = threadSense.get();
8     while (childCount.get() > 0) {};
9     childCount.set(children);
10    if (parent != null) {
11        parent.childDone();
12        while (sense != mySense) {};
13    } else {
14        sense = !sense;
15    }
16    threadSense.set(!mySense);
17 }
18 public void childDone() {
19     childCount.getAndDecrement();
20 }

```

# Barriers

- Dissemination Barriers
- Tournament Barriers
- Butterfly Barriers

...

- [Paper] Algorithms for Scalable Synchronization on Shared Memory Multiprocessors:  
<http://web.mit.edu/6.173/www/currentsemester/readings/R06-scalable-synchronization-1991.pdf>

# Case Study: Sum

- We have a large number of thread (think millions)
  - Each thread has an integer
  - The goal is to compute sum of all integers
- 
- Barriers!