



SIMON FRASER UNIVERSITY  
ENGAGING THE WORLD

# CMPT 431 Distributed Systems

Fall 2019

## Global State & Snapshot Recording

<https://www.cs.sfu.ca/~keval/teaching/cmpt431/fall19/>

Instructor: Keval Vora

# Reading

- [DC] Chapter 4
  - Upto 4.3



# Introduction

- Record the global state of a distributed system on-the-fly
- No global shared memory
- No global clock
- Unpredictable message delays

# Consistent Global State

- Global state is a collection of the local states of all processes, and the states of all channels


$$GS = \{ \cup_i LS_i, \cup_{i,j} SC_{ij} \}$$

- Global state GS is a consistent global state iff:

$$C1: \text{send}(m_{ij}) \in LS_i \Rightarrow m_{ij} \in SC_{ij} \oplus \text{rec}(m_{ij}) \in LS_j$$


$$C2: \text{send}(m_{ij}) \notin LS_i \Rightarrow m_{ij} \notin SC_{ij} \wedge \text{rec}(m_{ij}) \notin LS_j$$

# Consistent Global State



# Consistent Global State

Sum: \$800



# Recording a Global State

- How to distinguish between the messages to be recorded in the snapshot from those not to be recorded?
  - Any message that is sent by a process before recording its snapshot, must be recorded in the global snapshot (from C1)
  - Any message that is sent by a process after recording its snapshot, must not be recorded in the global snapshot (from C2)
- How to determine the instant when a process should take its snapshot?
  - Process  $p_j$  must record its snapshot before processing a message  $m_{ij}$  that was sent by process  $p_i$  after recording its snapshot

# Chandy-Lamport Algorithm

- Distributed algorithm to record global snapshot
- **Marker**: Control message to separate messages that should be included in the snapshot
- After a process records its snapshot, it sends a marker to all outgoing channels before sending other messages
- A process must record its snapshot no later than when it receives a marker on any of its incoming channels



# Chandy-Lamport Algorithm

- Initiated by any process by running ‘Marker Sending Rule’

*Marker sending rule* for process  $p_i$

- (1) Process  $p_i$  records its state.
- (2) For each outgoing channel C on which a marker has not been sent,  $p_i$  sends a marker along C before  $p_i$  sends further messages along C.

*Marker receiving rule* for process  $p_j$

On receiving a marker along channel C:

**if**  $p_j$  has not recorded its state **then**

Record the state of C as the empty set

Execute the “marker sending rule”

**else**


Record the state of C as the set of messages received along C after  $p_j$ 's state was recorded and before  $p_j$  received the marker along C

# Chandy-Lamport Algorithm: Correctness

- How to reason about correctness?
  - Does the algorithm satisfy C1 and C2?
- No message sent after the marker on that channel is recorded in the channel state (assumption: FIFO channel)
  - C2 is satisfied
- When a process  $p_j$  receives message  $m_{ij}$  that precedes the marker on channel  $C_{ij}$  :
  - If  $p_j$  has not taken its snapshot yet, it includes  $m_{ij}$  in its recorded snapshot
  - Otherwise, it records  $m_{ij}$  in the state of the channel  $C_{ij}$
  - C1 is satisfied

# Chandy-Lamport Algorithm

- 1<sup>st</sup> example:
  - $A = \$550$ ,  $B = \$170$
  - $C_{12} = \$0$ ,  $C_{21} = \$80$
- 2<sup>nd</sup> example:
  - $A = \$600$ ,  $B = \$120$
  - $C_{12} = \$0$ ,  $C_{21} = \$80$



Any issues?

# Chandy-Lamport Algorithm

- The recorded global state may not correspond to any of the global states that occurred during computation
- Reason: process can change its state **asynchronously** before the markers it sent are received by other sites and the other sites record their states
- The recorded state is a valid state that could happen
  - It retains “stable” properties (e.g., sum = \$800, no deadlock, etc.)
- Useful in detecting stable properties